

The slide features decorative watercolor splashes in the corners. The top-left corner has a cluster of green, blue, and purple splashes. The top-right corner has a few small orange and blue splashes. The bottom-left corner has several orange and red splashes. The bottom-right corner has a large, detailed splash with purple, blue, and green hues, including some fine lines and dots.

Intro Data Analysis in R

Lateef

University of Gloucestershire

Learning objectives:

- Understand and apply fundamental R tools to **read, clean** and **explore** datasets.
- **Create** and **interpret** basic data visualization using R tools.

Contents:

- Data Preprocessing
- Data Manipulation
- Data Visualization

Preliminaries:

Assumptions

- ✓ R and RStudio (or an alternative R-friendly interface) is installed on each student's computer.
- ➔ Students have a basic understanding of R, including familiarity with simple commands and data structures.
- 💡 Nevertheless, all R commands will be introduced in a clear and simplified manner to enhance the understanding of all learners.

Starter

✓ What is R? ✓ Packages installation:

- R needs packages to work. There are many available packages to use in R. E.g, *tidyverse*, *rio*, *lubridate*, *dplyr*, etc.
- To install a package use `install.package("package_name")`

✓ Libraries:

- Use to load **packages** into working file. E.g.
`library(package_name)`

```
1 library(tidyverse) # Load tidyverse into work session
2 library(gt)
3 library(readxl)
```

Data Preprocessing

In-built data

- There are several data available in R. This can be access via:

```
1 data() # Show all inbuilt data in R
```

Data loading

There are several ways of loading data in R, depending on the type of dataset.

```
1 data1 <- read_csv("dt_examp1.csv") # load csv file
2 data2 <- read_excel("dt_example.xlsx") # load excel file
3 data3 <- data(mpg) # load in-built data
```

Data exploring

After loading a dataset, it's important to explore and examine the dataset to gain a better understanding before proceeding with analysis.

```
1 str(data1)      # display the structure of the data
```

```
spc_tbl_ [1,000 × 15] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ TransactionID : chr [1:1000] "TX00001" "TX00002" "TX00003" "TX00004" ...
 $ BuyerName     : chr [1:1000] "Total Nigeria" "Eterna Plc" "Oando Plc" "Eterna Plc" ...
 $ FuelType      : chr [1:1000] "PMS" "AGO" "DPK" "DPK" ...
 $ SalesLiters   : num [1:1000] 43485 29355 22492 36538 24847 ...
 $ PricePerLitre : num [1:1000] 645 700 645 610 610 700 530 700 530 645 ...
 $ SalesDate     : Date[1:1000], format: "2025-02-28" "2025-01-09" ...
 $ Terminal      : chr [1:1000] "Port Harcourt Terminal" "Lagos Jetty" "Lekki Port" "Lagos Jetty" ...
 $ Operator      : chr [1:1000] "Shift C" "Shift B" "Shift B" "Shift B" ...
 $ VehicleType   : chr [1:1000] "Trailer" "Trailer" "Container" "Container" ...
 $ PaymentMethod : chr [1:1000] "Bank Transfer" "Cash" "Bank Transfer" "Bank Transfer" ...
 $ DeliveryStatus: chr [1:1000] "Cancelled" "Pending" "Cancelled" "Pending" ...
 $ Region        : chr [1:1000] "South-West" "North-East" "South-West" "South-East" ...
 $ InvoiceNumber  : chr [1:1000] "INV000001" "INV000002" "INV000003" "INV000004" ...
 $ BatchID       : chr [1:1000] "BATCH918" "BATCH345" "BATCH826" "BATCH262" ...
 $ TotalSales    : num [1:1000] 28047528 20548570 14507443 22287954 15156536 ...
 - attr(*, "spec")=
 .. cols(
 ..   TransactionID = col_character(),
 ..   BuyerName = col_character(),
 ..   _ = _
 .. )
```

Data exploring

```
1 dim(data1) # display the dimension of the data
```

```
[1] 1000 15
```

```
1 head(data1) # display the top 6 row of the data
```

```
# A tibble: 6 × 15
```

	TransactionID	BuyerName	FuelType	SalesLiters	PricePerLitre	SalesDate	Terminal
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<date>	<chr>
1	TX00001	Total Ni...	PMS	43485.	645	2025-02-28	Port Ha...
2	TX00002	Eterna P...	AGO	29355.	700	2025-01-09	Lagos J...
3	TX00003	Oando Plc	DPK	22492.	645	2025-03-22	Lekki P...
4	TX00004	Eterna P...	DPK	36538.	610	2025-04-13	Lagos J...
5	TX00005	Forte Oil	DPK	24847.	610	2025-03-10	Lekki P...
6	TX00006	Forte Oil	LPG	43798.	700	2025-01-17	Lekki P...

```
# i 8 more variables: Operator <chr>, VehicleType <chr>, PaymentMethod <chr>,
```

```
# DeliveryStatus <chr>, Region <chr>, InvoiceNumber <chr>, BatchID <chr>,
```

```
# TotalSales <dbl>
```


Data exploring

```
1 glimpse(data1) # display structural overview of the data
```

Rows: 1,000

Columns: 15

```
$ TransactionID <chr> "TX00001", "TX00002", "TX00003", "TX00004", "TX00005", ...
$ BuyerName    <chr> "Total Nigeria", "Eterna Plc", "Oando Plc", "Eterna Plc...
$ FuelType     <chr> "PMS", "AGO", "DPK", "DPK", "DPK", "LPG", "DPK", "AGO",...
$ SalesLiters  <dbl> 43484.54, 29355.10, 22492.16, 36537.63, 24846.78, 43797...
$ PricePerLitre <dbl> 645, 700, 645, 610, 610, 700, 530, 700, 530, 645, 700, ...
$ SalesDate    <date> 2025-02-28, 2025-01-09, 2025-03-22, 2025-04-13, 2025-0...
$ Terminal     <chr> "Port Harcourt Terminal", "Lagos Jetty", "Lekki Port", ...
$ Operator     <chr> "Shift C", "Shift B", "Shift B", "Shift B", "Shift D", ...
$ VehicleType  <chr> "Trailer", "Trailer", "Container", "Container", "Contai...
$ PaymentMethod <chr> "Bank Transfer", "Cash", "Bank Transfer", "Bank Transfe...
$ DeliveryStatus <chr> "Cancelled", "Pending", "Cancelled", "Pending", "Delive...
$ Region       <chr> "South-West", "North-East", "South-West", "South-East",...
$ InvoiceNumber <chr> "INV000001", "INV000002", "INV000003", "INV000004", "IN...
$ BatchID      <chr> "BATCH918", "BATCH345", "BATCH826", "BATCH262", "BATCH4...
$ TotalSales   <dbl> 28047528, 20548570, 14507443, 22287954, 15156536, 30658...
```

```
1 View(data1) # open the entire dataset in a new spreadsheet
```

Data Manipulation

Data subsetting

To filter our data, we first load the necessary package for data manipulation (dplyr)

```
1 filtered_data <- data1 %>% filter(SalesLiters > 30000)
2 filtered_data
```

```
# A tibble: 420 × 15
```

	TransactionID	BuyerName	FuelType	SalesLiters	PricePerLitre	SalesDate
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<date>
1	TX00001	Total Nigeria	PMS	43485.	645	2025-02-28
2	TX00004	Eterna Plc	DPK	36538.	610	2025-04-13
3	TX00006	Forte Oil	LPG	43798.	700	2025-01-17
4	TX00007	NNPC Retail	DPK	45134.	530	2025-02-24
5	TX00010	Eterna Plc	AGO	30025.	645	2025-02-24
6	TX00011	Ardova Plc	PMS	45706.	700	2025-06-17
7	TX00013	Oando Plc	PMS	31525.	645	2025-04-12
8	TX00014	NNPC Retail	PMS	31946.	610	2025-06-14
9	TX00015	Eterna Plc	AGO	36923.	645	2025-04-28
10	TX00017	Conoil	LPG	36075.	530	2025-03-09

```
# i 410 more rows
```

```
# i 9 more variables: Terminal <chr>, Operator <chr>, VehicleType <chr>,
```

```
#   PaymentMethod <chr>, DeliveryStatus <chr>, Region <chr>,
```

```
#   InvoiceNumber <chr>, BatchID <chr>, TotalSales <dbl>
```

Data group by

Observe that the filtered data contains multiple categories of fuel type. To gain deeper insights, we will group the data by fuel type and then summarize key statistics for each group.

```
1 filtered_data %>%  
2   group_by(FuelType) %>%  
3     summarise( mean(SalesLiters),  
4                 median(SalesLiters),  
5                 sd(SalesLiters))
```

A tibble: 4 × 4

	FuelType	`mean(SalesLiters)`	`median(SalesLiters)`	`sd(SalesLiters)`
	<chr>	<dbl>	<dbl>	<dbl>
1	AGO	39717.	40330.	5854.
2	DPK	40141.	40057.	5710.
3	LPG	39535.	39642.	6034.
4	PMS	38842.	38066.	5927.

Data cleaning

When preparing data for analysis, it is often necessary to modify the structure of your dataset by adding new columns, removing irrelevant ones, or renaming existing columns for clarity. To add a column, we use the **mutate()** function in *dplyr* as follows:

```
1 data1_more <- data1 %>%
2   mutate(Expenditure = SalesLiters * PricePerLitre) # Creating a new column
3
4 data1_h1 <- data1_more %>%
5   mutate(Price_Category = ifelse(PricePerLitre > 600, "High", "Low")) # Creating another column
```

```
1 data1_h1 <- data1 %>%
2   mutate(Expenditure = SalesLiters * PricePerLitre) %>%
3   mutate(Price_Category = ifelse(PricePerLitre > 600, "High", "Low"))
4
5 dim(data1_h1)
```

```
[1] 1000  17
```



Removing column

Similarly, columns can be removed by using **select()** function from *dplyr*.

```
1 data1_less <- data1 %>%  
2   select(-c(Region, InvoiceNumber)) # Removing column 'Region' and 'InvoiceNumber' from the data  
3  
4 dim(data1_less)
```

```
[1] 1000  13
```

Checking for missing values (NA)

Real-world datasets often contain missing values, which can negatively impact the quality of our analysis. It is good practice to identify and handle them before proceeding. In R, the `is.na()` function is commonly used to detect missing values.

```
1 x <- c(1, 2, NA, 4)
2 is.na(x)
```

```
[1] FALSE FALSE TRUE FALSE
```

...

```
1 sum(is.na(data1)) # Counting total missing values
```

```
[1] 0
```

```
1 colSums(is.na(data1)) # Finding missing values by column
```

TransactionID	BuyerName	FuelType	SalesLiters	PricePerLitre
0	0	0	0	0
SalesDate	Terminal	Operator	VehicleType	PaymentMethod
0	0	0	0	0
DeliveryStatus	Region	InvoiceNumber	BatchID	TotalSales
0	0	0	0	0

```
1 data1_clean <- na.omit(data1) # Removing NA in data
```



Replacing NA

Missing values can also be replaced with **mean**, **median**, **mode** or a **fixed value** in the column for completeness.

```
1 data1 <- data1 %>%  
2   mutate(SalesLiters = ifelse(is.na(SalesLiters),  
3                               mean(SalesLiters, na.rm = TRUE),  
4                               SalesLiters))
```

Summary of Filtered Data

```
1 filtered_data %>%  
2   group_by(Region) %>%  
3   summarise(  
4     Total_Liters = sum(SalesLiters),  
5     Total_Revenue = sum(TotalSales)  
6   )
```

A tibble: 6 × 3

	Region <chr>	Total_Liters <dbl>	Total_Revenue <dbl>
1	North-Central	2689212.	1676389685.
2	North-East	2462199.	1516967741.
3	North-West	3186724.	1975223068.
4	South-East	2576748.	1587752157.
5	South-South	3097156.	1929140745.
6	South-West	2599955	1608376159.

Data visualization

ggplot

ggplot is a powerful and flexible R package for data visualization, based on the **Grammar of Graphics**. It allows users to create complex, multi-layered plots in a structured and consistent way. Its basic syntax is


```
ggplot(data, aes(x = variable1, y = variable2)) + geom_style()
```

```
1 library(ggplot2) # loading the necessary package
```

Histogram

A histogram is a graphical representation of the distribution of a numeric variable. It divides the data into intervals (bins) and shows the count (or frequency) of values within each bin. For instance, Understanding the shape of the data (e.g., normal, skewed), Detecting outliers, spread, and central tendency.

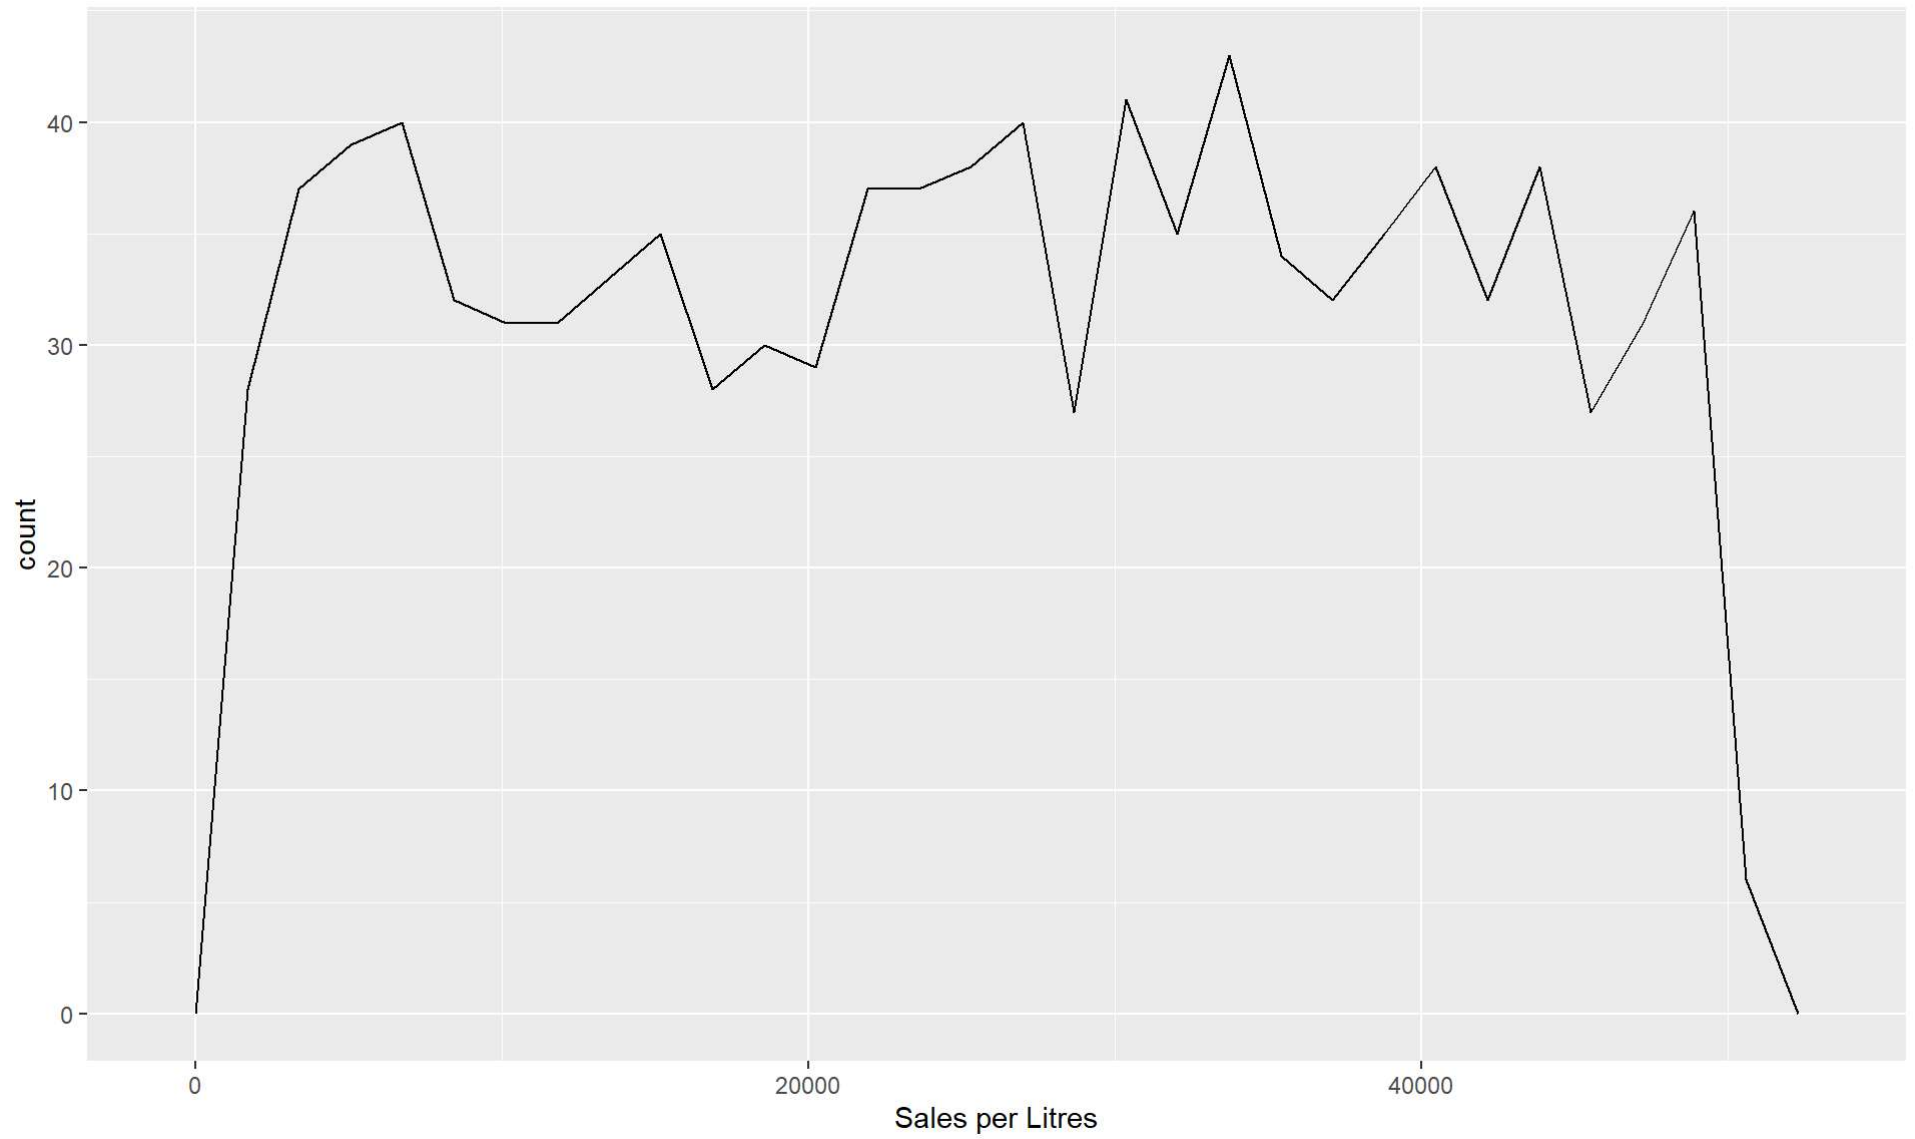
```
1 ggplot(data1, aes(x = SalesLiters)) +  
2   geom_histogram()           #Plotting a histogram
```



Frequency Polygon

A frequency polygon is a line graph that connects the midpoints of the top of each histogram bin. It shows the distribution shape more smoothly and is useful for comparing multiple distributions. It shows trends or patterns more clearly than histograms.


```
1 ggplot(data1, aes(x = SalesLiters)) +  
2   geom_freqpoly() +  
3   labs(x = "Sales per Litres")    # Plotting a frequency polygon
```



Histogram + Frequency Polygon

Typically, it is good to combine both a histogram and a frequency polygon to provide a more comprehensive view of the data distribution. Using the histogram to show actual counts per bin and the frequency polygon to highlight the overall shape or trend of the distribution.

```
1 ggplot(data1, aes(x = SalesLiters)) +  
2   geom_histogram() +  
3   geom_freqpoly() +  
4   labs(x = "Sales per Litres") +  
5   labs(y = "Frequency")
```





Bar Plots

A bar plot is used to display the frequency or value of categorical variables. Each bar represents a category, and the height of the bar corresponds to its value or count. It is typically used for:

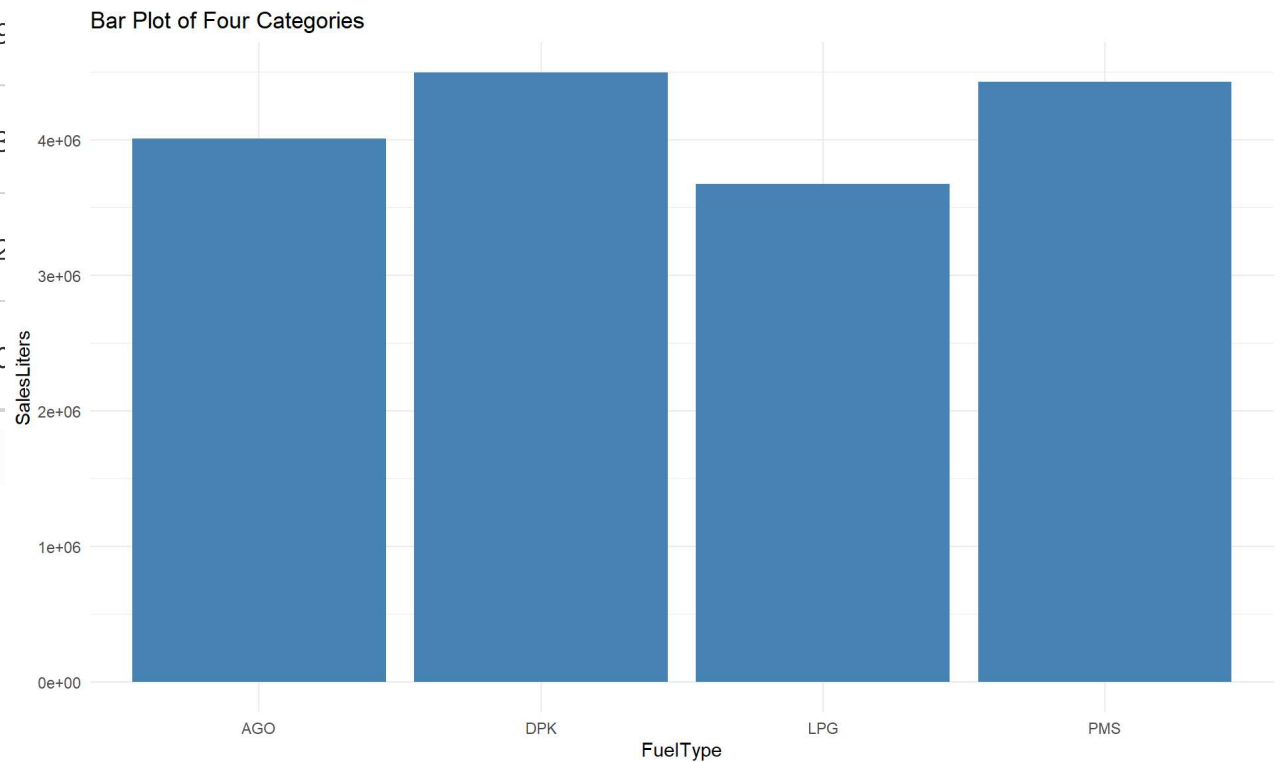
- Compare values of categories.
- Show the number of observations per category (e.g., fuel types, regions)
- Visualize grouped data.

Syntax : `ggplot(data, aes(x = Category, y = Value)) + geom_bar(stat = "identity")` # for pre-summarized values

Filtered data

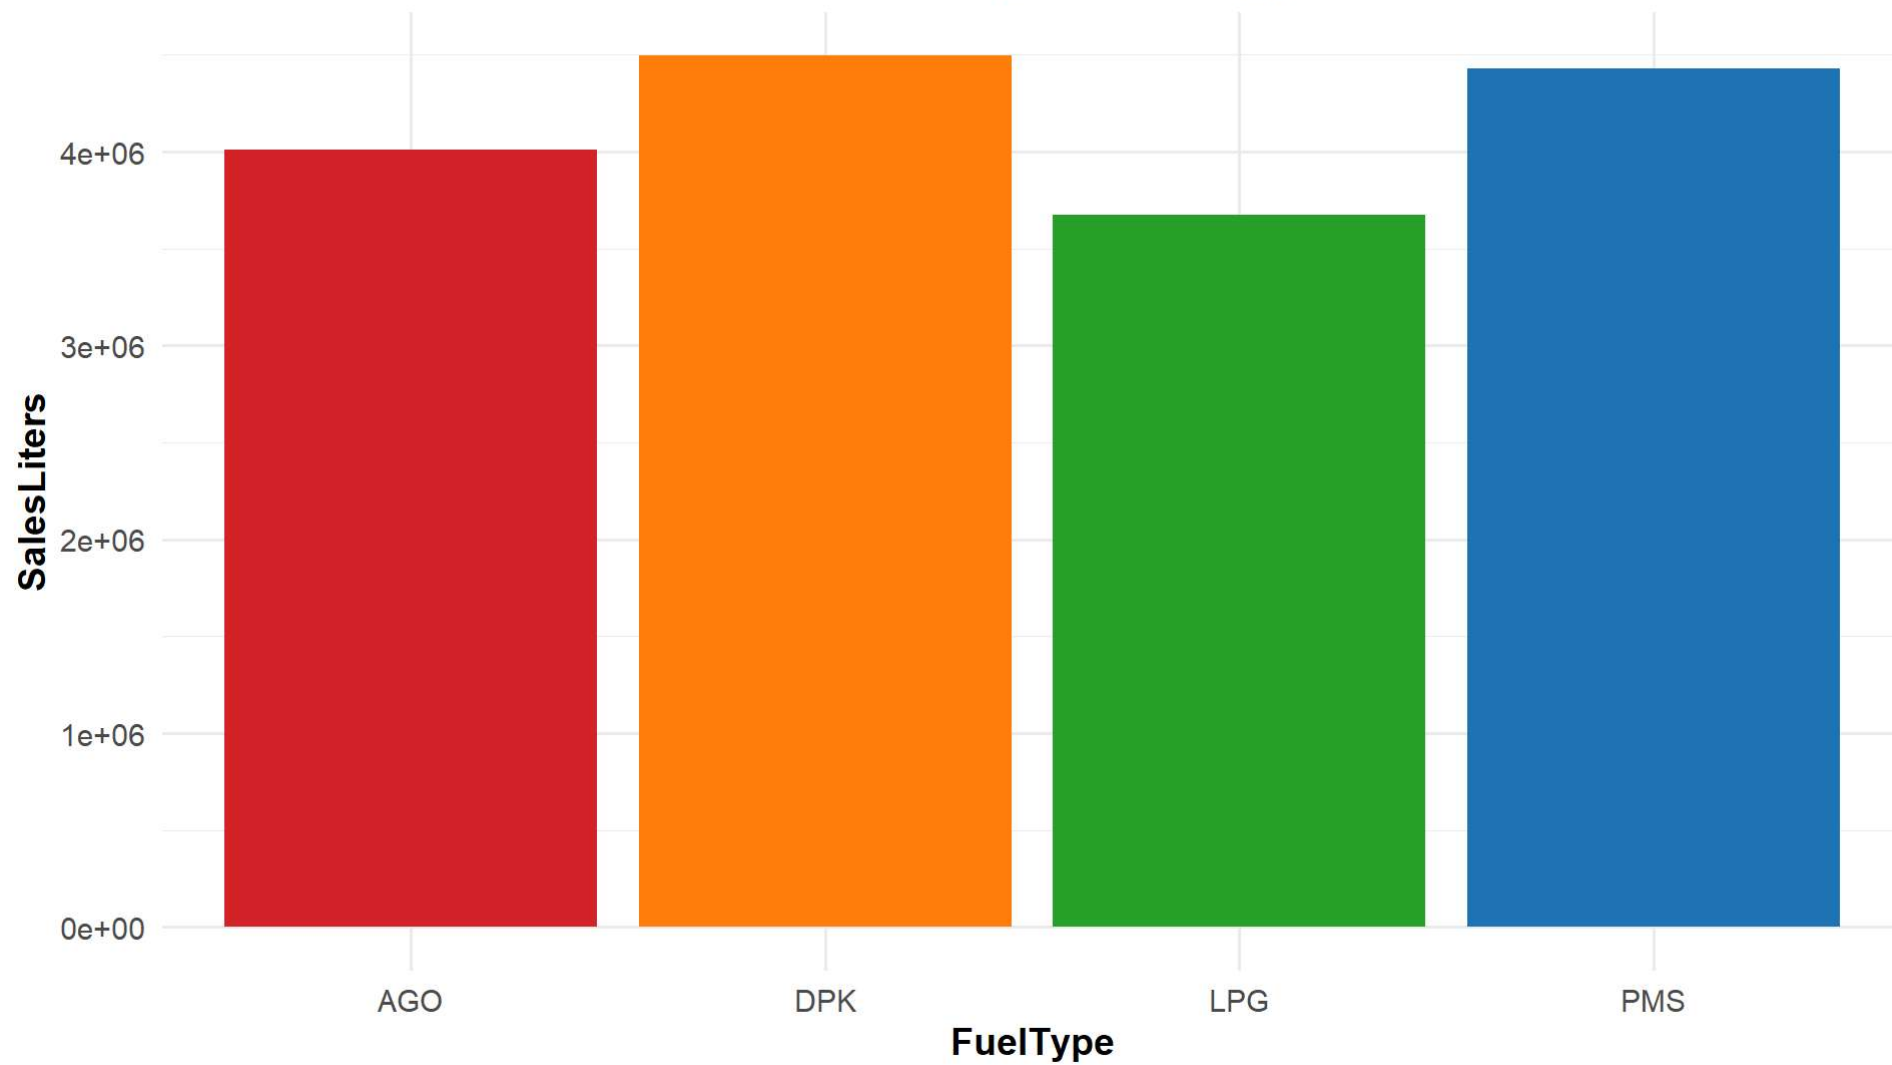
TransactionID	BuyerName	FuelType	SalesLiters
TX00001	Total Nigeria	PMS	4348
TX00004	Eterna Plc	DPK	3653
TX00006	Forte Oil	LPG	4379
TX00007	NNPC Retail	DPK	4513
TX00010	Eterna Plc	AGO	3002
TX00011	Ardova Plc	PMS	4570

```
1 ggplot(filtered_data, aes(x = FuelType, y = SalesLiters))
2   geom_bar(stat = "identity", fill = "steelblue") +
3   labs(title = "Bar Plot of Four Categories",
4         x = "FuelType",
5         y = "SalesLiters") +
6   theme_minimal()
```



```
1 ggplot(filtered_data, aes(x = FuelType, y = SalesLiters, fill = FuelType)) +  
2   geom_bar(stat = "identity") +  
3   labs(  
4     title = "SalesLitres per FuelTypes",  
5     x = "FuelType",  
6     y = "SalesLiters") +  
7   scale_fill_manual(values = c("PMS" = "#1f77b4",    # Blue  
8                                "DPK" = "#ff7f0e",    # Orange  
9                                "LPG" = "#2ca02c",    # Green  
10                               "AGO" = "#d62728")) + # Red  
11 theme_minimal(base_size = 14) +  
12   theme(  
13     plot.title = element_text(face = "bold", size = 20, hjust = 0.5),  
14     axis.title = element_text(face = "bold"),  
15     legend.position = "none"  
16   )
```

SalesLitres per FuelTypes




Line Plots

Line Plots

A line plot displays data points connected by straight lines. It is commonly used to visualize trends over time or continuous variables. E.g, Tracking monthly or daily sales, Visualizing temperature or price changes over time.

```
1 filtered_data %>%
2   mutate(Month = format(SalesDate, "%Y-%m")) %>%
3   group_by(Month) %>%
4   summarise(Monthly_Sales = sum(TotalSales)) %>%
5   ggplot(aes(x = Month, y = Monthly_Sales)) +
6   geom_line(group = 1, color = "steelblue") +
7   theme_minimal() +
8   labs(title = "Monthly Sales Trend", x = "Month", y = "Total Sales") +
9   theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Monthly Sales Trend



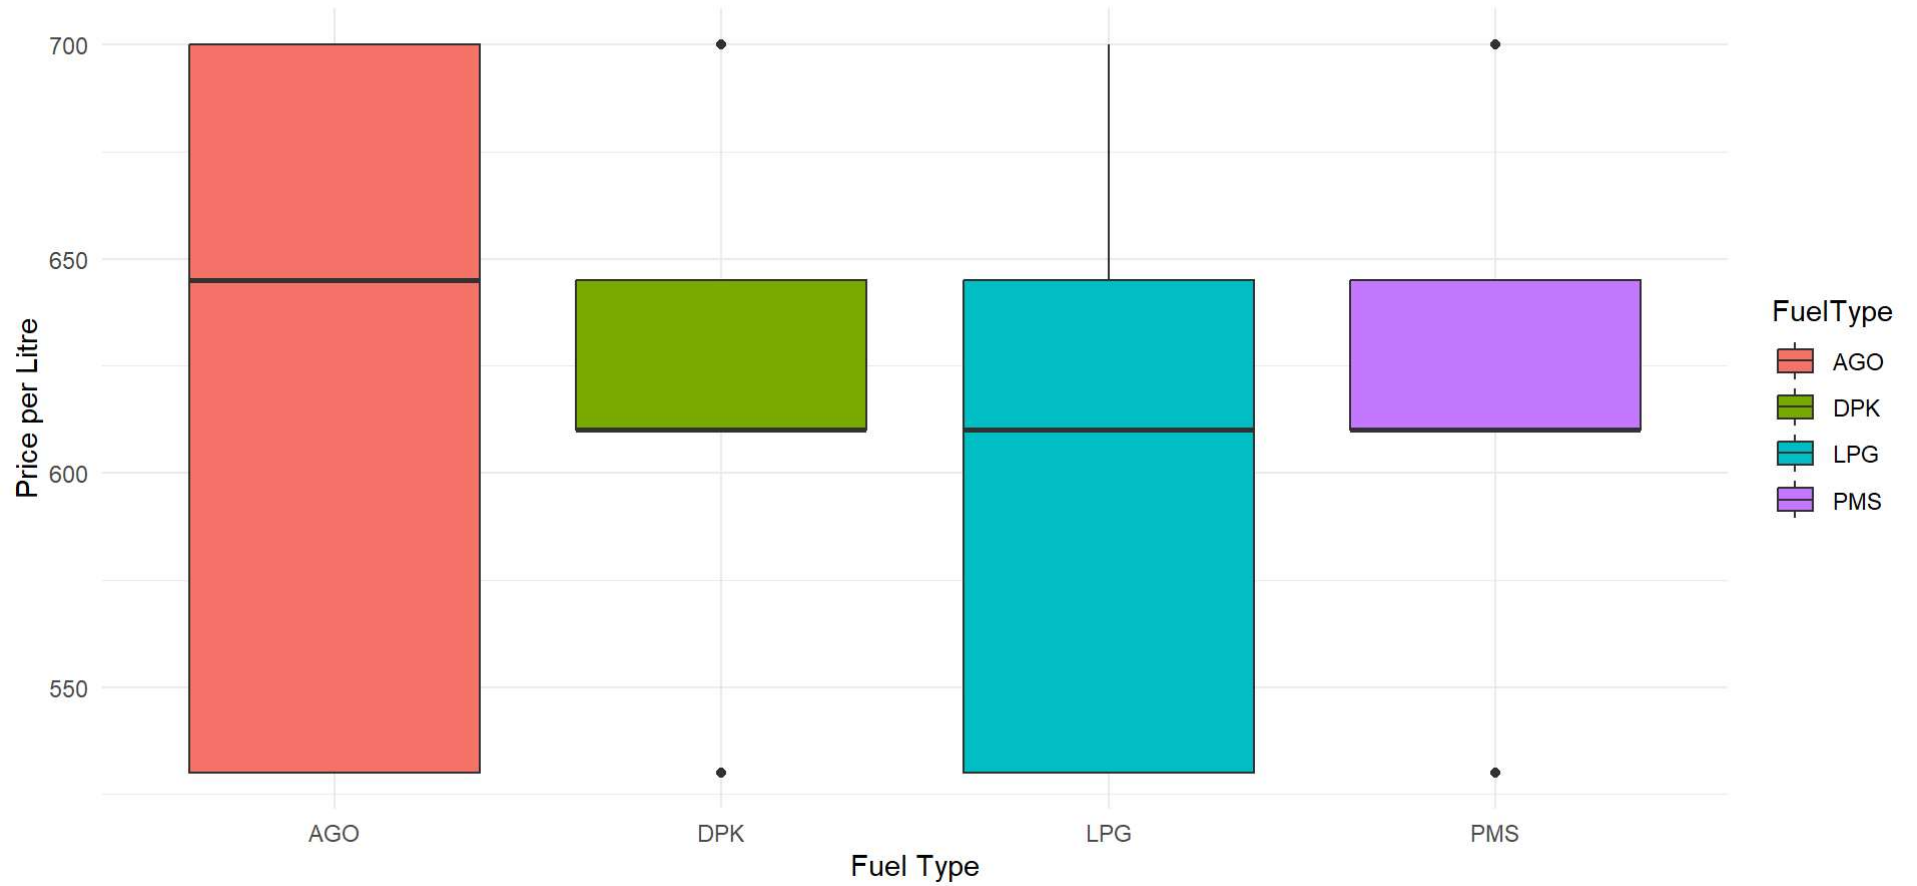
Box Plots

Box Plots

A box plot (also called a box-and-whisker plot) summarizes the distribution of a numeric variable using five statistics: Minimum, 1st quartile (Q1), Median, 3rd quartile (Q3), and Maximum.

```
1 ggplot(filtered_data, aes(x = FuelType, y = PricePerLitre, fill = FuelType)) +  
2   geom_boxplot() +  
3   labs(title = "Price Distribution by Fuel Type", x = "Fuel Type", y = "Price per  
4   theme_minimal() +  
5   theme(plot.title = element_text(face = "bold", size = 18, hjust = 0.5))
```

Price Distribution by Fuel Type



Summary

! Important

In this class, we have covered the following topics:

- Basic of data preprocessing: `str()`, `head()`, `dim()`, `glimpse()`, `View()`
- Basic of data manipulation using *dplyr*(): `filter()`, `group_by()`, `summarise()`, `mutate()`, `select()`, `is.na()`
- Basic of data visualization: `histogram()`, `freqplot()`, `barplot()`, `lineplot()`, `boxplot()`, customize plot.
- Rmarkdown: Use to prepare reports. **This note is prepared using Rmarkdown**